

```
from google.colab import drive
drive.mount('/content/gdrive', force_remount=True)
```

Mounted at /content/gdrive

```
import cv2
import torch
import numpy as np
import pandas as pd
from tensorflow.keras.applications.resnet_v2 import ResNet50V2, preprocess_input
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from tqdm import tqdm
```

```
#locate the file path
im_size = 224
train_x = []
train_y= []
val_x = []
val_y= []
dict_label={}
import os
trainfolder_list = os.listdir("/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/train")
valfolder_list = os.listdir("/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/val")
i=0
for folder in trainfolder_list:
    dict_label[folder] = i
    i+=1
```

```
#read training data
for folder in tqdm(trainfolder_list):
    for filename in os.listdir('/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/train/'+folder):
        img = cv2.resize(cv2.imread('/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/train/'+folder+'/'+filename,cv2.IMREAD_
img_array = preprocess_input(np.expand_dims(np.array(img[...,:-1].astype(np.float32)).copy(), axis=0))
train_x.append(img_array.reshape(3,im_size,im_size))
train_y.append(dict_label[folder])
```

100%|██████████| 10/10 [02:03<00:00, 12.33s/it]

```
#read validation data
for folder in tqdm(valfolder_list):
    for filename in os.listdir('/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/val/'+folder):
        img = cv2.resize(cv2.imread('/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/val/'+folder+'/'+filename,cv2.IMREAD_COLOR),
                           (im_size,im_size))
        img_array = preprocess_input(np.expand_dims(np.array(img[...::-1].astype(np.float32)).copy(), axis=0))
        val_x.append(img_array.reshape(3,im_size,im_size))
        val_y.append(dict_label[folder])
```


100%|██████████| 10/10 [00:03<00:00, 2.93it/s]

```
train_x, test_x, train_y, test_y = train_test_split(train_x, train_y, test_size=0.2, random_state=42)
```

```
#just checking the data
plt.imshow(train_x[150].reshape(224,224,3))
plt.show()
print(train_y[150])
print(dict_label)

print("images-size:", train_x[0].shape)
print(len(train_y))
print(len(val_y))
print(len(test_y))
print(val_y)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



```
from torch.utils.data import Dataset, DataLoader, ConcatDataset
batchSize = 32
dataset=list(zip(train_x, train_y))
trainloader = DataLoader(dataset, batch_size = batchSize, shuffle=True)
```

```
validdataset=list(zip(val_x, val_y))
valiloader = DataLoader(validdataset, batch_size = batchSize, shuffle=True)
```

```
testset=list(zip(test_x, test_y))
testloader = DataLoader(testset, batch_size = batchSize, shuffle=True)
```

```
from torchvision import *
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
net = models.resnet18(pretrained=True)
net = net.cuda() if device else net
net
```

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
```

```

        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(layer2): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)

```

```

import torch.nn as nn
import torch.optim as optim
criterion = nn.CrossEntropyLoss()

```

```
optimizer = optim.SGD(net.parameters(), lr=0.0001, momentum=0.9)

def accuracy(out, labels):
    _,pred = torch.max(out, dim=1)
    return torch.sum(pred==labels).item()
num_ftrs = net.fc.in_features
net.fc = nn.Linear(num_ftrs, 128)
net.fc = net.fc.cuda() if torch.cuda.is_available() else net.fc
```

```
n_epochs = 200
print_every = 10
valid_loss_min = np.Inf
val_loss = []
val_acc = []
train_loss = []
train_acc = []
total_step = len(trainloader)
for epoch in range(1, n_epochs+1):
    running_loss = 0.0
    correct = 0
    total=0
    print(f'Epoch {epoch}\n')
    for batch_idx, (data_, target_) in enumerate(trainloader):
        data_, target_ = data_.to(device), target_.to(device)
        optimizer.zero_grad()
        outputs = net(data_)
        loss = criterion(outputs, target_)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
        _,pred = torch.max(outputs, dim=1)
        correct += torch.sum(pred==target_).item()
        total += target_.size(0)
        if (batch_idx) % 20 == 0:
            print ('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}'
                  .format(epoch, n_epochs, batch_idx, total_step, loss.item()))
    train_acc.append(100 * correct / total)
    train_loss.append(running_loss/total_step)
    print(f'\ntrain-loss: {np.mean(train_loss):.4f}, train-acc: {(100 * correct/total):.4f}')
    batch_loss = 0
    total_t=0
    correct_t=0
    with torch.no_grad():
        for batch_idx, (data_, target_) in enumerate(validloader):
            data_, target_ = data_.to(device), target_.to(device)
            outputs = net(data_)
            loss = criterion(outputs, target_)
            val_loss.append(loss.item())
            _,pred = torch.max(outputs, dim=1)
            correct_t += torch.sum(pred==target_).item()
            total_t += target_.size(0)
        val_acc.append(100 * correct_t / total_t)
        valid_loss_min = min(valid_loss_min, np.mean(val_loss))
        if np.mean(val_loss) < valid_loss_min:
            save_path = './mnist_model.pth'
            torch.save(net.state_dict(), save_path)
            print('Model saved')
```

```

with torch.no_grad():
    net.eval()
    for data_t, target_t in (testloader):
        data_t, target_t = data_t.to(device), target_t.to(device)
        outputs_t = net(data_t)
        loss_t = criterion(outputs_t, target_t)
        batch_loss += loss_t.item()
        _,pred_t = torch.max(outputs_t, dim=1)
        correct_t += torch.sum(pred_t==target_t).item()
        total_t += target_t.size(0)
    val_acc.append(100 * correct_t/total_t)
    val_loss.append(batch_loss/len(testloader))
    network_learned = batch_loss < valid_loss_min
    print(f'validation loss: {np.mean(val_loss):.4f}, validation acc: {(100 * correct_t/total_t):.4f}\n')

if network_learned:
    valid_loss_min = batch_loss
    torch.save(net.state_dict(), 'resnet.pt')
    print('Improvement-Detected, save-model')

net.train()

```

Epoch 1

```

Epoch [1/200], Step [0/312], Loss: 5.2945
Epoch [1/200], Step [20/312], Loss: 5.2177
Epoch [1/200], Step [40/312], Loss: 5.0701
Epoch [1/200], Step [60/312], Loss: 5.1168
Epoch [1/200], Step [80/312], Loss: 5.1485
Epoch [1/200], Step [100/312], Loss: 5.1567
Epoch [1/200], Step [120/312], Loss: 5.0575
Epoch [1/200], Step [140/312], Loss: 5.1182
Epoch [1/200], Step [160/312], Loss: 5.3125
Epoch [1/200], Step [180/312], Loss: 5.0191
Epoch [1/200], Step [200/312], Loss: 5.1787
Epoch [1/200], Step [220/312], Loss: 5.0047
Epoch [1/200], Step [240/312], Loss: 5.0725
Epoch [1/200], Step [260/312], Loss: 5.0997
Epoch [1/200], Step [280/312], Loss: 5.2080
Epoch [1/200], Step [300/312], Loss: 4.9834

```

```

train-loss: 5.1032, train-acc: 0.2206
validation loss: 5.0082, validation acc: 0.4813

```

Improvement-Detected, save-model

Epoch 2

```
Epoch [2/200], Step [0/312], Loss: 5.0730
Epoch [2/200], Step [20/312], Loss: 4.8301
Epoch [2/200], Step [40/312], Loss: 4.8943
Epoch [2/200], Step [60/312], Loss: 4.9916
Epoch [2/200], Step [80/312], Loss: 5.0898
Epoch [2/200], Step [100/312], Loss: 4.9159
Epoch [2/200], Step [120/312], Loss: 4.9604
Epoch [2/200], Step [140/312], Loss: 5.0017
Epoch [2/200], Step [160/312], Loss: 4.8191
Epoch [2/200], Step [180/312], Loss: 4.9980
Epoch [2/200], Step [200/312], Loss: 4.9156
Epoch [2/200], Step [220/312], Loss: 4.8459
Epoch [2/200], Step [240/312], Loss: 4.7125
Epoch [2/200], Step [260/312], Loss: 4.8236
Epoch [2/200], Step [280/312], Loss: 4.7720
Epoch [2/200], Step [300/312], Loss: 4.7647
```

train-loss: 4.9994, train-acc: 1.9055

validation loss: 4.9092, validation acc: 4.7734


Improvement-Detected, save-model

Epoch 3

```
Epoch [3/200], Step [0/312], Loss: 4.8815
Epoch [3/200], Step [20/312], Loss: 4.6747
Epoch [3/200], Step [40/312], Loss: 4.8179
Epoch [3/200], Step [60/312], Loss: 4.5073
Epoch [3/200], Step [80/312], Loss: 4.7630
Epoch [3/200], Step [100/312], Loss: 4.8134
Epoch [3/200], Step [120/312], Loss: 4.5645
Epoch [3/200], Step [140/312], Loss: 4.8515
Epoch [3/200], Step [160/312], Loss: 4.7804
Epoch [3/200], Step [180/312], Loss: 4.6670
```

```
fig = plt.figure(figsize=(10,5))
plt.title("Train-Validation Accuracy")
plt.plot(train_acc, label='train')
plt.plot(val_acc, label='validation')
plt.xlabel('num_epochs', fontsize=12)
plt.ylabel('accuracy', fontsize=12)
plt.legend(loc='best')
```

<matplotlib.legend.Legend at 0x7f1f6c903710>



```
correct = 0
total = 0
pred_list=[]
true_list=[]
# since we're not training, we don't need to calculate the gradients for our outputs
net.eval()
with torch.no_grad():
    for data in valiloaders:
        images, labels = data
        images, labels = images.cuda(), labels.cuda()
        # calculate outputs by running images through the network
        outputs = net(images)
        # the class with the highest energy is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        pred_list.append(predicted.cpu().numpy())
        true_list.append(labels.cpu().numpy())

print(f'Accuracy of the network on the 10000 test images: {100 * correct // total} %')
```

Accuracy of the network on the 10000 test images: 34 %


```

pred_list=list(np.concatenate(pred_list).flat)
true_list=list(np.concatenate(true_list).flat)
print(pred_list)
print(true_list)

```

```

[6, 8, 1, 9, 8, 4, 3, 0, 5, 2, 2, 1, 3, 7, 4, 8, 9, 7, 9, 2, 7, 7, 2, 4, 5, 1, 6, 3, 8, 9, 9, 7, 0, 4, 1, 1, 6, 8, 4, 1, 2, 7, 3,
[6, 8, 9, 9, 2, 6, 3, 0, 5, 2, 0, 4, 9, 3, 4, 0, 4, 1, 0, 5, 1, 3, 1, 4, 9, 1, 4, 1, 5, 4, 4, 7, 9, 0, 1, 1, 6, 5, 1, 7, 3, 0, 5,

```

```

from sklearn.metrics import confusion_matrix, classification_report
array=confusion_matrix(true_list, pred_list)
print(array)

```

```

[[ 7  7  1  3  4  3  3  7  5 10]
 [ 2 15  3  4  9  0  3  8  1  5]
 [ 1  0 32  1  3  4  1  1  5  2]
 [ 1  1  2 14  3  3 11  8  2  5]
 [ 2  5  2  5 11  4  9  4  1  7]
 [ 1  2  6  4  3 16  1  3  8  6]
 [ 3  3  2  4  4  2 18  6  5  3]
 [ 0  8  1  3  7  0  5 20  1  5]
 [ 1  1  4  7  4  7  2  0 21  3]
 [ 7  5  4  5  2  2  3  1  3 18]]

```

```

#dummy classifier
from sklearn.dummy import DummyClassifier
X = true_list
y = true_list
dummy_clf = DummyClassifier(strategy="most_frequent") #or most_frequent,stratified
dummy_clf.fit(X, y)
print("dummy classifier result\n")
#print(dummy_clf.score(X, y))
x_dummy=dummy_clf.predict(X)
print(classification_report(true_list, x_dummy))

```

dummy classifier result

| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.10 | 1.00 | 0.18 | 50 |
| 1 | 0.00 | 0.00 | 0.00 | 50 |

| | | | | |
|---|------|------|------|----|
| 2 | 0.00 | 0.00 | 0.00 | 50 |
| 3 | 0.00 | 0.00 | 0.00 | 50 |
| 4 | 0.00 | 0.00 | 0.00 | 50 |
| 5 | 0.00 | 0.00 | 0.00 | 50 |
| 6 | 0.00 | 0.00 | 0.00 | 50 |
| 7 | 0.00 | 0.00 | 0.00 | 50 |
| 8 | 0.00 | 0.00 | 0.00 | 50 |
| 9 | 0.00 | 0.00 | 0.00 | 50 |

| | | | | |
|--------------|------|------|------|-----|
| accuracy | | | 0.10 | 500 |
| macro avg | 0.01 | 0.10 | 0.02 | 500 |
| weighted avg | 0.01 | 0.10 | 0.02 | 500 |

```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))

```




```

#plot the confusion_matrix, and classification_report
import seaborn as sn
print("confusion matrix: ")
df_cm = pd.DataFrame(array, range(10), range(10))
sn.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
plt.xlabel('predicted label', fontsize=18)
plt.ylabel('true label', fontsize=16)
plt.show()
print("\n\nclassification report:\n\n"+classification_report(true_list, pred_list))

```

confusion matrix:



classification report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.28 | 0.14 | 0.19 | 50 |
| 1 | 0.32 | 0.30 | 0.31 | 50 |
| 2 | 0.56 | 0.64 | 0.60 | 50 |
| 3 | 0.28 | 0.28 | 0.28 | 50 |
| 4 | 0.22 | 0.22 | 0.22 | 50 |
| 5 | 0.39 | 0.32 | 0.35 | 50 |
| 6 | 0.32 | 0.36 | 0.34 | 50 |
| 7 | 0.34 | 0.40 | 0.37 | 50 |
| 8 | 0.40 | 0.42 | 0.41 | 50 |
| 9 | 0.28 | 0.36 | 0.32 | 50 |
| accuracy | | | 0.34 | 500 |
| macro avg | 0.34 | 0.34 | 0.34 | 500 |
| weighted avg | 0.34 | 0.34 | 0.34 | 500 |

