

```
from google.colab import drive
drive.mount('/content/gdrive', force_remount=True)
```

Mounted at /content/gdrive

```
import cv2
import torch
import numpy as np
import pandas as pd
from tensorflow.keras.applications.resnet_v2 import ResNet50V2, preprocess_input
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from tqdm import tqdm
```

```
#locate the file path
im_size = 224
train_x = []
train_y= []
val_x = []
val_y= []
dict_label={}
import os
trainfolder_list = os.listdir("/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/train")
valfolder_list = os.listdir("/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/val")
i=0
for folder in trainfolder_list:
    dict_label[folder] = i
    i+=1
```

```
#read training data
for folder in tqdm(trainfolder_list):
    for filename in os.listdir('/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/train/'+folder):
        img = cv2.resize(cv2.imread('/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/train/'+folder+'/'+filename,cv2.IMREAD_
img_array = preprocess_input(np.expand_dims(np.array(img[...::-1].astype(np.float32)).copy(), axis=0))
train_x.append(img_array.reshape(3,im_size,im_size))
train_y.append(dict_label[folder])
```

100%|██████████| 10/10 [02:03<00:00, 12.33s/it]

```
#read validation data
for folder in tqdm(valfolder_list):
    for filename in os.listdir('/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/val/'+folder):
        img = cv2.resize(cv2.imread('/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/val/'+folder+'/'+filename,cv2.IMREAD_COLOR),
                           (im_size,im_size))
        img_array = preprocess_input(np.expand_dims(np.array(img[...,:-1].astype(np.float32)).copy(), axis=0))
        val_x.append(img_array.reshape(3,im_size,im_size))
        val_y.append(dict_label[folder])
```

100%|██████████| 10/10 [00:03<00:00, 2.93it/s]

```
train_x, test_x, train_y, test_y = train_test_split(train_x, train_y, test_size=0.2, random_state=42)
```

```
#just checking the data
plt.imshow(train_x[150].reshape(224,224,3))
plt.show()
print(train_y[150])
print(dict_label)

print("images-size:", train_x[0].shape)
print(len(train_y))
print(len(val_y))
print(len(test_y))
print(val_y)
```



```

        (relu): ReLU(inplace=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(layer2): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)

```

```

import torch.nn as nn
import torch.optim as optim
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.0001, momentum=0.9)

```

```

def accuracy(out, labels):
    _,pred = torch.max(out, dim=1)
    return torch.sum(pred==labels).item()
num_fts = net.fc.in_features
net.fc = nn.Linear(num_fts, 128)
net.fc = net.fc.cuda() if torch.cuda.is_available() else net.fc

```

```

n_epochs = 50
print_every = 10
valid_loss_min = np.Inf
val_loss = []
val_acc = []
train_loss = []
train_acc = []
total_step = len(trainloader)
for epoch in range(1, n_epochs+1):
    running_loss = 0.0
    correct = 0
    total=0
    print(f'Epoch {epoch}\n')
    for batch_idx, (data_, target_) in enumerate(trainloader):
        data_, target_ = data_.to(device), target_.to(device)
        optimizer.zero_grad() # zero the parameter gradients
        outputs = net(data_) # forward
        loss = criterion(outputs, target_)
        loss.backward() #backward
        optimizer.step() #optimize
        running_loss += loss.item()
        _,pred = torch.max(outputs, dim=1) # result of train
        correct += torch.sum(pred==target_).item()
        total += target_.size(0)
        if (batch_idx) % 20 == 0:
            print ('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}'
                  .format(epoch, n_epochs, batch_idx, total_step, loss.item()))
    train_acc.append(100 * correct / total)
    train_loss.append(running_loss/total_step)
    print(f'\ntrain-loss: {np.mean(train_loss):.4f}, train-acc: {(100 * correct/total):.4f}')
    batch_loss = 0
    total_t=0
    correct_t=0
    with torch.no_grad():

```

```

net.eval()
for data_t, target_t in (testloader):
    data_t, target_t = data_t.to(device), target_t.to(device)
    outputs_t = net(data_t)
    loss_t = criterion(outputs_t, target_t)
    batch_loss += loss_t.item()
    _,pred_t = torch.max(outputs_t, dim=1)    # result of validation
    correct_t += torch.sum(pred_t==target_t).item()
    total_t += target_t.size(0)
val_acc.append(100 * correct_t/total_t)
val_loss.append(batch_loss/len(testloader))
network_learned = batch_loss < valid_loss_min
print(f'validation loss: {np.mean(val_loss):.4f}, validation acc: {(100 * correct_t/total_t):.4f}\n')

if network_learned:
    valid_loss_min = batch_loss
    torch.save(net.state_dict(), 'resnet.pt')
    print('Improvement-Detected, save-model')

```

```
net.train()
```

Epoch 1

```

Epoch [1/50], Step [0/312], Loss: 5.2005
Epoch [1/50], Step [20/312], Loss: 5.0639
Epoch [1/50], Step [40/312], Loss: 5.2119
Epoch [1/50], Step [60/312], Loss: 5.0445
Epoch [1/50], Step [80/312], Loss: 4.8600
Epoch [1/50], Step [100/312], Loss: 5.0268
Epoch [1/50], Step [120/312], Loss: 5.1811
Epoch [1/50], Step [140/312], Loss: 5.0006
Epoch [1/50], Step [160/312], Loss: 4.9041
Epoch [1/50], Step [180/312], Loss: 5.1079
Epoch [1/50], Step [200/312], Loss: 4.9323
Epoch [1/50], Step [220/312], Loss: 4.8471
Epoch [1/50], Step [240/312], Loss: 5.0265
Epoch [1/50], Step [260/312], Loss: 5.0492
Epoch [1/50], Step [280/312], Loss: 5.0017
Epoch [1/50], Step [300/312], Loss: 5.1574

```

```

train-loss: 5.0641, train-acc: 0.2407
validation loss: 4.9505, validation acc: 0.9226

```

Improvement-Detected, save-model

Epoch 2

Epoch [2/50], Step [0/312], Loss: 4.8512  
Epoch [2/50], Step [20/312], Loss: 4.9664  
Epoch [2/50], Step [40/312], Loss: 5.1286  
Epoch [2/50], Step [60/312], Loss: 4.8805  
Epoch [2/50], Step [80/312], Loss: 4.7874  
Epoch [2/50], Step [100/312], Loss: 4.7970  
Epoch [2/50], Step [120/312], Loss: 4.9842  
Epoch [2/50], Step [140/312], Loss: 4.8065  
Epoch [2/50], Step [160/312], Loss: 4.9451  
Epoch [2/50], Step [180/312], Loss: 4.7618  
Epoch [2/50], Step [200/312], Loss: 4.8417  
Epoch [2/50], Step [220/312], Loss: 4.8299  
Epoch [2/50], Step [240/312], Loss: 4.6699  
Epoch [2/50], Step [260/312], Loss: 4.9697  
Epoch [2/50], Step [280/312], Loss: 4.6094  
Epoch [2/50], Step [300/312], Loss: 4.7161

train-loss: 4.9608, train-acc: 2.2365

validation loss: 4.8401, validation acc: 5.3751


Improvement-Detected, save-model

Epoch 3

Epoch [3/50], Step [0/312], Loss: 4.7666  
Epoch [3/50], Step [20/312], Loss: 4.8541  
Epoch [3/50], Step [40/312], Loss: 5.1191  
Epoch [3/50], Step [60/312], Loss: 4.5958  
Epoch [3/50], Step [80/312], Loss: 4.9548  
Epoch [3/50], Step [100/312], Loss: 4.7945  
Epoch [3/50], Step [120/312], Loss: 4.6657  
Epoch [3/50], Step [140/312], Loss: 4.5556  
Epoch [3/50], Step [160/312], Loss: 4.6119

```
fig = plt.figure(figsize=(10,5))
plt.title("Train-Validation Accuracy")
plt.plot(train_acc, label='train')
plt.plot(val_acc, label='validation')
plt.xlabel('num_epochs', fontsize=12)
plt.ylabel('accuracy', fontsize=12)
plt.legend(loc='best')
```

<matplotlib.legend.Legend at 0x7f1f579b7210>



```
correct = 0
total = 0
pred_list=[]
true_list=[]
# since we're not training, we don't need to calculate the gradients for our outputs
net.eval()
with torch.no_grad():
    for data in valiloaders:
        images, labels = data
        images, labels = images.cuda(), labels.cuda()
        # calculate outputs by running images through the network
        outputs = net(images)
        # the class with the highest energy is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        pred_list.append(predicted.cpu().numpy())
        true_list.append(labels.cpu().numpy())


print(f'Accuracy of the network on the test images: {100 * correct // total} %')
```

Accuracy of the network on the 10000 test images: 37 %



```
pred_list=list(np.concatenate(pred_list).flat)
true_list=list(np.concatenate(true_list).flat)
print(pred_list)
print(true_list)
```

```
[2, 5, 6, 9, 9, 3, 5, 8, 9, 6, 8, 4, 8, 0, 6, 7, 7, 3, 2, 8, 4, 6, 4, 0, 2, 5, 3, 7, 3, 8, 1, 2, 6, 1, 5, 4, 1, 9, 5, 2, 4, 7, 3,
[4, 5, 1, 1, 9, 6, 5, 9, 3, 6, 8, 4, 8, 0, 4, 3, 0, 6, 4, 8, 5, 1, 1, 0, 2, 5, 8, 7, 8, 8, 1, 2, 6, 7, 2, 4, 1, 8, 5, 2, 6, 1, 3,
```



```
from sklearn.metrics import confusion_matrix, classification_report
array=confusion_matrix(true_list, pred_list)
print(array)
```

```
[[10  1  5  1  1  1  2  8  8 13]
 [ 1 20  2  4  8  1  4  5  2  3]
 [ 0  2 39  1  2  2  0  1  3  0]
 [ 0  2  3 16  5  4  3  7  4  6]
 [ 0  6 10  2 13  0  5  1  6  7]
 [ 4  1  5  1  4 24  2  1  4  4]
 [ 0  3  0 10  5  4 14  6  5  3]
 [ 2 12  0  5  7  0  2 17  1  4]
 [ 0  1  3  6  3  7  4  1 20  5]
 [ 5  2  3  1  9  2  1  5  6 16]]
```

```
#dummy classifier
from sklearn.dummy import DummyClassifier
X = true_list
y = true_list
dummy_clf = DummyClassifier(strategy="most_frequent") #or most_frequent,stratified
dummy_clf.fit(X, y)
print("dummy classifier result\n")
#print(dummy_clf.score(X, y))
x_dummy=dummy_clf.predict(X)
print(classification_report(true_list, x_dummy))
```

dummy classifier result

	precision	recall	f1-score	support
0	0.10	1.00	0.18	50
1	0.00	0.00	0.00	50
2	0.00	0.00	0.00	50
3	0.00	0.00	0.00	50

4	0.00	0.00	0.00	50
5	0.00	0.00	0.00	50
6	0.00	0.00	0.00	50
7	0.00	0.00	0.00	50
8	0.00	0.00	0.00	50
9	0.00	0.00	0.00	50
accuracy			0.10	500
macro avg	0.01	0.10	0.02	500
weighted avg	0.01	0.10	0.02	500

```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are
  _warn_prf(average, modifier, msg_start, len(result))


```

```

#plot the confusion_matrix, and classification_report
import seaborn as sn
print("confusion matrix: ")
df_cm = pd.DataFrame(array, range(10), range(10))
sn.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
plt.xlabel('predicted label', fontsize=18)
plt.ylabel('true label', fontsize=16)
plt.show()
print("\n\nclassification report:\n\n"+classification_report(true_list, pred_list))

```

confusion matrix:



classification report:

	precision	recall	f1-score	support
0	0.45	0.20	0.28	50
1	0.40	0.40	0.40	50
2	0.56	0.78	0.65	50
3	0.34	0.32	0.33	50
4	0.23	0.26	0.24	50
5	0.53	0.48	0.51	50
6	0.38	0.28	0.32	50
7	0.33	0.34	0.33	50
8	0.34	0.40	0.37	50
9	0.26	0.32	0.29	50
accuracy			0.38	500
macro avg	0.38	0.38	0.37	500
weighted avg	0.38	0.38	0.37	500

