

```
from google.colab import drive
drive.mount('/content/gdrive', force_remount=True)
```

Mounted at /content/gdrive

```
import cv2
import numpy as np
import pandas as pd
from tensorflow.keras.applications.resnet_v2 import ResNet50V2, preprocess_input
import matplotlib.pyplot as plt
```

```
from sklearn.metrics import classification_report
from tqdm import tqdm
```

```
im_size = 224
train_x = []
train_y= []
val_x = []
val_y= []
dict_label={}
import os
trainfolder_list = os.listdir("/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/train")
valfolder_list = os.listdir("/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/val")
i=0
for folder in trainfolder_list:
    dict_label[folder] = i
    i+=1
```

```
for folder in tqdm(trainfolder_list):
    for filename in os.listdir('/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/train/'+folder):
        img = cv2.resize(cv2.imread('/content/gdrive/MyDrive/FindCareer/FellowshipAI/imagewoof-320/train/'+folder+'/'+filename, cv2.IMREAD_
        img_array = preprocess_input(np.expand_dims(np.array(img[...::-1].astype(np.float32)).copy(), axis=0))
        train_x.append(img_array.reshape(3, im_size, im_size))
        train_y.append(dict_label[folder])
```

100%|██████████| 10/10 [01:48<00:00, 10.88s/it]


```
import torch
import torchvision.models as models
model = models.resnet18(pretrained = True)
model

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
# Freeze the parameters
for param in model.parameters():
    param.requires_grad = False
```

```
#Classifier architecture to put on top of resnet18
from torch import nn
from collections import OrderedDict
fc = nn.Sequential(OrderedDict([
    ('fc1', nn.Linear(512,100)),
    ('relu', nn.ReLU()),
    ('fc2', nn.Linear(100,10)),
    ('output', nn.LogSoftmax(dim=1))
]))

model.fc = fc
```

```
#shifting model to gpu
model.to(device)
model
```

```
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
```

```

        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Sequential(
  (fc1): Linear(in_features=512, out_features=100, bias=True)
  (relu): ReLU()
  (fc2): Linear(in_features=100, out_features=10, bias=True)
  (output): LogSoftmax(dim=1)
)
)

```

```

from torch.utils.data import Dataset, DataLoader, ConcatDataset
dataset=list(zip(train_x, train_y))

```

```
dataloader = DataLoader(dataset, batch_size = 64, shuffle=True)

validataset=list(zip(val_x, val_y))
valiloader = DataLoader(validataset, batch_size = 64, shuffle=False)
```

```
from torch import optim
def train(model, trainloader, criterion, optimizer, epochs = 5):
    train_loss =[]
    validate_loss =[]
    for e in range(epochs):
        running_loss =0
        for images, labels in trainloader:
            inputs, labels = images.to(device), labels.to(device)
            optimizer.zero_grad()
            img = model(inputs)

            loss = criterion(img, labels)
            running_loss+=loss
            loss.backward()
            optimizer.step()
        print("Epoch : {}/{}..".format(e+1,epochs),
              "Training Loss: {:.6f}".format(running_loss/len(train_y)))
        train_loss.append(running_loss)
```

```
epochs = 7
model.train()
optimizer = optim.Adam(model.fc.parameters(), lr=0.001)
criterion = nn.NLLLoss()
train(model,dataloader,criterion, optimizer, epochs)
```

```
Epoch : 1/7.. Training Loss: 0.033367
Epoch : 2/7.. Training Loss: 0.031000
Epoch : 3/7.. Training Loss: 0.030367
Epoch : 4/7.. Training Loss: 0.029788
Epoch : 5/7.. Training Loss: 0.029524
Epoch : 6/7.. Training Loss: 0.029287
Epoch : 7/7.. Training Loss: 0.028981
```




```
x_dummy=dummy_clf.predict(X)
print(classification_report(val_y, x_dummy))
```

dummy classifier result

	precision	recall	f1-score	support
0	0.07	0.08	0.07	50
1	0.08	0.10	0.09	50
2	0.22	0.26	0.24	50
3	0.10	0.08	0.09	50
4	0.06	0.04	0.05	50
5	0.04	0.04	0.04	50
6	0.09	0.10	0.09	50
7	0.00	0.00	0.00	50
8	0.07	0.08	0.07	50
9	0.10	0.10	0.10	50
accuracy			0.09	500
macro avg	0.08	0.09	0.08	500
weighted avg	0.08	0.09	0.08	500

```
#plot the confusion_matrix, and classification_report
import seaborn as sn
df_cm = pd.DataFrame(array, range(10), range(10))
sn.heatmap(df_cm, annot=True, annot_kws={"size": 16}, fmt='g')
plt.xlabel('predicted label', fontsize=18)
plt.ylabel('true label', fontsize=16)
plt.show()
print("\n\nclassification report:\n\n"+classification_report(val_y, pred_list))
```



classification report:

	precision	recall	f1-score	support
0	0.50	0.10	0.17	50
1	0.30	0.38	0.34	50
2	0.66	0.50	0.57	50
3	0.32	0.14	0.19	50
4	0.22	0.16	0.19	50
5	0.36	0.18	0.24	50
6	0.31	0.44	0.37	50
7	0.20	0.44	0.28	50
8	0.28	0.32	0.30	50
9	0.24	0.34	0.28	50
accuracy			0.30	500
macro avg	0.34	0.30	0.29	500
weighted avg	0.34	0.30	0.29	500

